

1. Intro & Design Twitter

Twitter Clone - Step-by-Step ERD & Database Design Notes

This document explains the **entity-relationship diagram (ERD)** design for a **Twitter-like social media app**. The goal is to build a clear understanding of database modeling by starting from **feature requirements**, then building **intuitively** and **iteratively** step-by-step.

Purpose of this Document

- Designed for **SDE-1 interview prep** and **self-learning**
 - Follows an **intuition-first** approach before implementation
 - Breaks down complex concepts into digestible steps
 - SQL/PostgreSQL-centric thinking (constraints, structure, integrity)
-

Step 1: Requirements Breakdown

Let's begin by **understanding the application's core features**:

Feature Set:

- Users can **sign up**
- Users can **post tweets** (text + optional media)
- Users can **follow other users**
- Users can **like tweets**
- Users can **comment** on tweets
- Users can **subscribe to premium** (e.g., blue checkmark)

From this, we extract the **main entities** and **actions**:

Core Entities (Tables):

1. Users
 2. Tweets
 3. Media (attached to tweets)
 4. Follows (many-to-many)
 5. Likes
 6. Comments
 7. Subscriptions
-

Step 2: Understanding Primary Keys & Business Logic

What is a Primary Key?

A **primary key (PK)** is a column (or set of columns) that **uniquely identifies each row** in a table.

Properties of a Primary Key:

- Must be **unique** across the table
 - Cannot be **NULL**
 - Should be **immutable** (doesn't change over time)
 - Optimized for fast lookups (often indexed by default)
-

What is Business Logic in a Database?

Business logic refers to **real-world rules and decisions** related to your application's features or workflows.

Examples in our app:

- `username` is chosen by users
- `email` is required for login/communication
- `plan_type` decides user privileges

🚫 Why Not Use Business Logic as a Primary Key?

1. User-controlled and changeable

- A user may want to change their `username` or `email`
- Changing PKs causes complex cascading updates

2. Validation-heavy

- Requires frequent uniqueness checks at the app level
- Email format can vary or cause issues (e.g., uppercase vs lowercase)

3. Security implications

- Exposing sensitive or meaningful data (like email) via URLs or APIs is dangerous

4. Lack of global uniqueness

- Business values often assume uniqueness *per context* (e.g., `username` in one org)

✅ Preferred Approach: Use a Surrogate Key

- Use a **UUID** (universally unique identifier) or **CUID** as `id`
- Auto-generated by backend or database

Example in PostgreSQL:

```
CREATE TABLE users (  
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
  username VARCHAR(255) UNIQUE NOT NULL,  
  email VARCHAR(255) UNIQUE NOT NULL,  
  bio TEXT,  
  created_at TIMESTAMP DEFAULT NOW()  
);
```

Here, `id` is the primary key — not `username` or `email`.

Step 3: Start with the Most Fundamental Entity - Users

Every action in our system starts with a user — posting tweets, liking, following, etc.

Key Points:

- `id` is the primary key (UUID)
- `username` and `email` should be unique
- Always include `created_at`

We'll keep the design simple and scalable — more fields can be added later based on use case.

Step 4: Moving Forward - Planning Relationships

Before jumping into the rest of the tables, let's build strong **intuition** for:

- **One-to-One:** e.g., A user has one subscription
- **One-to-Many:** e.g., A user can post many tweets
- **Many-to-Many:** e.g., Users follow each other

We'll explore these as we move entity-by-entity. No code for ERD generation yet — we'll focus on logic first.

Stay tuned for:

- Tweets
- Media
- Comments
- Likes
- Follows
- Subscriptions

Each will be broken down step-by-step with **justification** before implementing any code.

What's Next

Next, we will:

- Go entity by entity
- Understand **real-world behavior** first
- Justify each relationship (1:1, 1:N, N:N)
- Plan for constraints, normalization, scalability

Step 5: Tweets Entity - Thinking in Relationships

Real-World Behavior

- A user can post **many tweets**
- A tweet can only be created by **one user**

This is a textbook **One-to-Many (1:N)** relationship:

- One `user` → many `tweets`
- Each `tweet` → one `user`

Database Design Thoughts:

- `tweets` table should have a foreign key `user_id` referring to `users(id)`
- This allows us to query: *"show me all tweets by this user"*
- Every tweet should be timestamped (`created_at`)
- Optional: add `updated_at` if we allow tweet edits

Suggested Fields in `tweets` :

Column	Type	Constraints
id	UUID	PRIMARY KEY
user_id	UUID	FOREIGN KEY → users(id)
content	TEXT	NOT NULL
created_at	TIMESTAMP	DEFAULT NOW()

Column	Type	Constraints
updated_at	TIMESTAMP	NULLABLE (if edits allowed)

Example Query:

```
CREATE TABLE tweets (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  user_id UUID NOT NULL REFERENCES users(id),
  content TEXT NOT NULL,
  created_at TIMESTAMP DEFAULT NOW(),
  updated_at TIMESTAMP
);
```

Step 6: Media Entity - Handling Tweet Attachments

Real-World Behavior

- A tweet can contain **zero, one, or multiple media files** (images, videos, gifs)
- Each media item **belongs to exactly one tweet**

This leads to a **One-to-Many (1:N)** relationship:

- One `tweet` → many `media`
- Each `media` → one `tweet`

Database Design Thoughts:

- `media` table stores URL and type of media (image, video, etc.)
- Foreign key `tweet_id` refers to the `tweets(id)`
- Include timestamps for tracking uploads
- Optional: a `position` field if media has ordering (e.g., first image, second image)

Suggested Fields in `media` :

Column	Type	Constraints
id	UUID	PRIMARY KEY
tweet_id	UUID	FOREIGN KEY → tweets(id)
url	TEXT	NOT NULL
media_type	TEXT	e.g., 'image', 'video', 'gif'
created_at	TIMESTAMP	DEFAULT NOW()

Example Query:

```
CREATE TABLE media (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  tweet_id UUID NOT NULL REFERENCES tweets(id),
  url TEXT NOT NULL,
  media_type TEXT NOT NULL,
  created_at TIMESTAMP DEFAULT NOW()
);
```

In the next section, we'll move to the **comments** entity and understand the dual relationship it has with both users and tweets.

Step 7: Comments Entity - Dual Relationships

In any social media platform, **comments** are crucial for user interaction and engagement.

We'll now explore how to model the **comments** entity logically and relationally in our database.

Real-World Behavior

- A **user** can post **many comments**
- A **tweet** can receive **many comments**
- A **comment** must be associated with **exactly one user** (who authored it)
- A **comment** must also belong to **exactly one tweet** (the thread it's under)

💬 Two Distinct Relationships

This makes `comments` a **junction entity** that connects:

1. `users` → `comments` (1:N)
2. `tweets` → `comments` (1:N)

+ Optional Add-ons

- We might want to allow nested replies (comments on comments), but we'll leave that out for now for simplicity.

🔧 Database Design Thoughts

The `comments` table should:

- Have a unique ID (UUID)
- Store comment content (`TEXT`)
- Store `user_id` as a foreign key to reference the author
- Store `tweet_id` as a foreign key to reference the tweet
- Include a `created_at` timestamp

⚙️ Suggested Fields in `comments` :

Column	Type	Constraints
id	UUID	PRIMARY KEY
user_id	UUID	FOREIGN KEY → users(id)
tweet_id	UUID	FOREIGN KEY → tweets(id)
content	TEXT	NOT NULL
created_at	TIMESTAMP	DEFAULT NOW()

🧠 Why This Design Makes Sense

- Each `comment` must be traced back to both:
 - Who created it? → `user_id`

- On which tweet? → `tweet_id`
- This setup allows queries like:
 - "Get all comments on this tweet"
 - "Get all comments made by this user"
- Ensures **data normalization** and **referential integrity**

SQL Snippet (for reference)

```
CREATE TABLE comments (  
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
  user_id UUID NOT NULL REFERENCES users(id),  
  tweet_id UUID NOT NULL REFERENCES tweets(id),  
  content TEXT NOT NULL,  
  created_at TIMESTAMP DEFAULT NOW()  
);
```

In the next section, we'll tackle the **likes** entity, which connects users and tweets in a many-to-many relationship.

We'll explore how to structure a clean, efficient join table for it.

💖 Step 8: Likes Entity - Many-to-Many Relationships

Likes are a fundamental part of user interaction in any social media app. In a Twitter-like system, users can like multiple tweets, and each tweet can be liked by multiple users.

This is a **classic many-to-many (M:N) relationship** between `users` and `tweets`.

💡 Real-World Behavior

- A **user** can like **many tweets**
- A **tweet** can be liked by **many users**
- Each **like** is specific to one `user` and one `tweet`

Thus, we need a **join table** to map this relationship.

Why a Separate Likes Table?

Many-to-many relationships should not be modeled by directly embedding arrays in relational DBs.

Instead, we use a join table — in this case, `likes`.

Each row in the `likes` table represents **one unique like action** performed by a user on a tweet.

Database Design Thoughts

The `likes` table should:

- Have its own `id` as primary key (UUID)
- Store `user_id` (who liked)
- Store `tweet_id` (which tweet was liked)
- Have a `created_at` timestamp

Optionally, you can enforce a unique constraint on (`user_id`, `tweet_id`) to prevent duplicate likes.

Suggested Fields in `likes` :

Column	Type	Constraints
<code>id</code>	UUID	PRIMARY KEY
<code>user_id</code>	UUID	FOREIGN KEY → <code>users(id)</code>
<code>tweet_id</code>	UUID	FOREIGN KEY → <code>tweets(id)</code>
<code>created_at</code>	TIMESTAMP	DEFAULT NOW()

Why This Design Makes Sense

- Prevents duplication of likes with a `(user_id, tweet_id)` uniqueness constraint
- Cleanly tracks who liked what and when
- Supports queries like:
 - "Get all tweets liked by a user"
 - "Count how many likes a tweet received"
 - "Check if a specific user liked a tweet"

SQL Snippet (for reference)

```
CREATE TABLE likes (  
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
  user_id UUID NOT NULL REFERENCES users(id),  
  tweet_id UUID NOT NULL REFERENCES tweets(id),  
  created_at TIMESTAMP DEFAULT NOW(),  
  CONSTRAINT unique_like UNIQUE (user_id, tweet_id)  
);
```

In the next section, we'll handle the **follows** relationship — a more complex **self-referencing many-to-many** relationship between users.

Step 9: Follows Entity - Self-Referencing Many-to-Many

Following relationships are at the heart of a social network like Twitter. This feature allows users to follow other users, resulting in a **many-to-many** relationship **between users themselves**.

Real-World Behavior

- A **user** can follow **many other users**
- A **user** can be followed by **many users**

This clearly represents a **self-referencing many-to-many relationship**.

How to Represent This in SQL

We can't use a direct relationship within the `users` table for this. Instead, we create a **join table**, typically named `follows`.

This table records each follow action:

- Who is the **follower**?
- Who is the **followee**?

Each row represents: *User A follows User B*

Database Design Thoughts

The `follows` table should:

- Have a unique `id` (UUID)
 - Have two foreign keys:
 - `follower_id` → `users(id)`
 - `followee_id` → `users(id)`
 - Include a `created_at` timestamp
 - Enforce a uniqueness constraint on `(follower_id, followee_id)`
-

Suggested Fields in `follows` :

Column	Type	Constraints
<code>id</code>	UUID	PRIMARY KEY
<code>follower_id</code>	UUID	FOREIGN KEY → <code>users(id)</code>
<code>followee_id</code>	UUID	FOREIGN KEY → <code>users(id)</code>
<code>created_at</code>	TIMESTAMP	DEFAULT NOW()

Constraints to Consider

- Prevent duplicate follow records with a unique pair constraint

- Prevent users from following themselves

Additional Constraints (Recommended)

```
CHECK (follower_id <> followee_id)
```

Why This Design Makes Sense

- Clearly models "who follows whom"
- Efficient for querying both followers and followees
- Prevents redundant or self-follow entries

Supports queries like:

- "List all followers of a user"
- "List everyone this user is following"
- "Count followers/following"

SQL Snippet (for reference)

```
CREATE TABLE follows (  
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
  follower_id UUID NOT NULL REFERENCES users(id),  
  followee_id UUID NOT NULL REFERENCES users(id),  
  created_at TIMESTAMP DEFAULT NOW(),  
  CONSTRAINT unique_follow UNIQUE (follower_id, followee_id),  
  CONSTRAINT no_self_follow CHECK (follower_id <> followee_id)  
);
```

Next up, we'll finalize with the **subscriptions** entity — representing the optional premium plan for users.

Step 10: Subscriptions Entity - One-to-One Relationship

In modern social media platforms, offering premium features (like a verified badge) is common. In our Twitter-like system, this is handled by **subscriptions**.

Each user can have **at most one active subscription**, making this a **one-to-one relationship** between `users` and `subscriptions`.

💡 Real-World Behavior

- A **user** can optionally subscribe to a premium plan
- Each **subscription** belongs to exactly **one user**
- A **user** can have **only one active subscription** at a time

This is modeled as:

- One `user` → one `subscription`
 - One `subscription` → one `user`
-

🔧 Database Design Thoughts

The `subscriptions` table should:

- Have a unique `id` (UUID)
 - Include a `user_id` field (foreign key and unique)
 - Store plan details (e.g., plan type)
 - Include start and end timestamps
 - Optional: `is_active` boolean flag
-

⚙️ Suggested Fields in `subscriptions` :

Column	Type	Constraints
<code>id</code>	UUID	PRIMARY KEY
<code>user_id</code>	UUID	FOREIGN KEY → <code>users(id)</code> , UNIQUE
<code>plan_type</code>	TEXT	e.g., 'basic', 'premium', 'verified'
<code>start_date</code>	DATE	NOT NULL

Column	Type	Constraints
end_date	DATE	NULLABLE (null means active)
created_at	TIMESTAMP	DEFAULT NOW()

Why This Design Makes Sense

- Enforces **one-to-one** via a **UNIQUE** constraint on **user_id**
- Can easily check if a user is subscribed
- Allows future enhancements:
 - Plan history table
 - Grace periods, renewals
 - Payment linkage

SQL Snippet (for reference)

```
CREATE TABLE subscriptions (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  user_id UUID UNIQUE NOT NULL REFERENCES users(id),
  plan_type TEXT NOT NULL,
  start_date DATE NOT NULL,
  end_date DATE,
  created_at TIMESTAMP DEFAULT NOW()
);
```

With this, we've completed the high-level database modeling of all core entities in our Twitter-like application. Ready to proceed with indexing, constraints, ERD generation, or additional features!

Step 10: Subscriptions Entity - One-to-One Relationship

In modern social media platforms, offering premium features (like a verified badge) is common. In our Twitter-like system, this is handled by **subscriptions**.

Each user can have **at most one active subscription**, making this a **one-to-one relationship** between `users` and `subscriptions`.

💡 Real-World Behavior

- A **user** can optionally subscribe to a premium plan
- Each **subscription** belongs to exactly **one user**
- A **user** can have **only one active subscription** at a time

This is modeled as:

- One `user` → one `subscription`
 - One `subscription` → one `user`
-

🔧 Database Design Thoughts

The `subscriptions` table should:

- Have a unique `id` (UUID)
 - Include a `user_id` field (foreign key and unique)
 - Store plan details (e.g., plan type)
 - Include start and end timestamps
 - Optional: `is_active` boolean flag
-

⚙️ Suggested Fields in `subscriptions` :

Column	Type	Constraints
<code>id</code>	UUID	PRIMARY KEY
<code>user_id</code>	UUID	FOREIGN KEY → <code>users(id)</code> , UNIQUE
<code>plan_type</code>	TEXT	e.g., 'basic', 'premium', 'verified'
<code>start_date</code>	DATE	NOT NULL

Column	Type	Constraints
end_date	DATE	NULLABLE (null means active)
created_at	TIMESTAMP	DEFAULT NOW()

Why This Design Makes Sense

- Enforces **one-to-one** via a **UNIQUE** constraint on **user_id**
- Can easily check if a user is subscribed
- Allows future enhancements:
 - Plan history table
 - Grace periods, renewals
 - Payment linkage

SQL Snippet (for reference)

```
CREATE TABLE subscriptions (  
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),  
  user_id UUID UNIQUE NOT NULL REFERENCES users(id),  
  plan_type TEXT NOT NULL,  
  start_date DATE NOT NULL,  
  end_date DATE,  
  created_at TIMESTAMP DEFAULT NOW()  
);
```

With this, we've completed the high-level database modeling of all core entities in our Twitter-like application. Ready to proceed with indexing, constraints, ERD generation, or additional features!



🧠 Deep-Dive Tasks for the Reader

◆ Task 1: Add `retweets` Feature

In Twitter, users can retweet other users' tweets.

Design a `retweets` table that:

- Tracks `user_id` (who retweeted)
- Tracks `original_tweet_id` (which tweet was retweeted)
- Includes `created_at` timestamp

Challenge: Can a retweet also have a comment (like quote-retweet)? Model that case.

◆ Task 2: Normalize Media Types

Instead of storing media types as free text, move them to a new table.

Steps:

- Create a `media_types` table (`id` , `name`)
- Link `media` to `media_types` via foreign key

This is useful for enums or controlled vocabularies.

◆ Task 3: Add Indexing Strategy

Interviewers love asking: "How would you improve performance?"

Try this:

- Identify which fields should be indexed.
 - Think about queries: search by `username` , `tweet_id` , etc.
 - Add indexes in SQL or write index plans.
-

◆ Task 4: Create an Analytics Table

For example: tracking how many tweets a user makes per day.

Your challenge:

- Design a `user_activity` table that stores:
 - `user_id`
 - `date`
 - `tweet_count`
 - `like_count`
 - Think: Should this be precomputed or live queried?
-

◆ Task 7: Design Query Challenges

Imagine you're the backend dev. Write raw SQL for the following:

1. Get all tweets by a user with their media.
 2. Count how many likes a tweet has.
 3. Find all users who follow each other mutually.
 4. List users with active subscriptions.
 5. Get the top 5 most commented tweets in the last week.
-

◆ Task 8: Extend for Notifications

Users get notified when someone likes, comments, or follows them.

Design a `notifications` table:

- Fields: `id`, `recipient_user_id`, `type`, `actor_user_id`, `related_tweet_id`, `created_at`, `read`
-